

Does Unemployment Risk Affect Business Cycle Dynamics?

By SEBASTIAN GRAVES*

In this paper, I show that the decline in consumption during unemployment depends on both liquid and illiquid wealth; that unemployment predicts illiquid asset withdrawal, primarily when households have few liquid assets; and that increased idiosyncratic unemployment risk leads to a rise in saving overall, but also to a decline in investment in illiquid assets. Motivated by these new findings, I embed endogenous unemployment risk in a two-asset heterogeneous-agent New Keynesian model. The model is consistent with the new evidence and suggests that aggregate shocks are amplified by a flight-to-liquidity when unemployment risk rises, particularly when monetary policy is constrained.

I. Introduction

Unemployment spells are the largest source of income risk that households face. Yet the majority of household wealth is held in illiquid assets, which are not well suited to smoothing consumption during unemployment. In this paper, I study the implications of these facts in a model with endogenous unemployment risk in which households trade both liquid and illiquid assets. The combination of these features provides a propagation mechanism for aggregate shocks, driven by a “flight-to-liquidity” that occurs when households face higher unemployment risk. It also suggests that an important role for unemployment insurance is its ability to dampen this amplification by lessening the cyclicalities of household income risk.

I begin by presenting new empirical evidence on the relationship between unemployment, the liquidity of asset holdings, and consumption. First, using data from the Consumer Expenditure Survey and the Panel Study of Income Dynamics, I show that the size of the decline in consumption during unemployment spells depends on both a household’s liquid and illiquid asset positions. In particular, the

* Federal Reserve Board. Email: sebastian.h.graves@frb.gov. I am very grateful to Simon Gilchrist, Mark Gertler, and Thomas Sargent for their advice and support throughout this project. I would like to thank Corina Boar, Jaroslav Borovička, Katarína Borovičková, Francesco Ferrante, William Gamber, Nils Gornemann, Ralph Luetticke (discussant), Virgiliu Midrigan, Gaston Navarro, Ayşegül Şahin (editor), and two anonymous referees for helpful comments, as well as seminar participants at NYU, the Federal Reserve Board, and Washington University in St. Louis. The views expressed in this paper are solely those of the author and should not be interpreted as reflecting the views of the Board of Governors of the Federal Reserve System or any other person associated with the Federal Reserve System.

consumption decline is smallest for households with significant liquid asset holdings, larger for households with only illiquid assets, and largest for households with few assets of either type.¹ This finding suggests that households principally use their liquid wealth to smooth consumption during unemployment, but are partially able to access their illiquid wealth if required. This is consistent with my second finding: I use data from the Survey of Consumer Finances to show that unemployment spells predict withdrawal from illiquid asset holdings, primarily when the unemployment spell is long or when the household has few liquid assets.

Taken together, the above evidence on consumption and illiquid asset adjustment during unemployment spells suggests that it is costly for households to use their illiquid wealth to smooth consumption during unemployment, and consequently that households will vary their liquid and illiquid asset holdings depending on the unemployment risk that they face. I use data from the Survey of Consumer Expectations to show that this is the case. I find that saving increases overall when idiosyncratic unemployment risk rises, but that investment in illiquid assets declines, consistent with a precautionary flight-to-liquidity mechanism.

Motivated by this empirical evidence, I study a heterogeneous-agent New Keynesian (HANK) model in which households trade both liquid bonds and illiquid capital, and are subject to endogenous unemployment risk due to search frictions in the labor market. First, I show that this model is consistent with all the above findings. I then study the response of the economy to aggregate shocks in order to answer the following questions: How does household demand for liquid and illiquid assets change when unemployment risk rises? Does this affect business cycle dynamics? Does unemployment insurance play an important role as an automatic stabilizer?

I find that the interaction of illiquid assets and endogenous unemployment risk provides an important propagation mechanism for aggregate shocks. Higher unemployment risk triggers a flight-to-liquidity: households increase their demand for liquid assets, as these are the best suited to smoothing consumption during unemployment spells. Conversely, demand for illiquid capital declines. In the presence of sticky prices, this decline in investment leads to lower output and higher unemployment, prompting a feedback loop between unemployment risk and aggregate demand. A key feature of the model is that income risk responds endogenously to aggregate shocks through changes in the unemployment rate. I use the Current Population Survey to show that this is confirmed in the data: the cyclicity of income risk is driven by changes in unemployment risk.²

¹In the terminology of Kaplan, Violante and Weidner (2014), the first group are non hand-to-mouth households, the second are the wealthy hand-to-mouth, and the third are the poor hand-to-mouth.

²Güvenen, Ozkan and Song (2014) use Social Security Administration data to show that the skewness of the income growth distribution is strongly pro-cyclical. The role of unemployment cannot be studied in their data, as it does not include a measure of time spent employed.

If there is no unemployment insurance, this propagation mechanism implies that the response of unemployment or output is around 30% larger than in a version of the model with no idiosyncratic unemployment risk. Unemployment insurance provides a source of consumption smoothing during unemployment spells, and consequently dampens the flight-to-liquidity and the feedback loop between unemployment risk and aggregate demand. Quantitatively, I find that unemployment insurance removes around half of the amplification that the flight-to-liquidity mechanism provides. Unemployment insurance is even more important when monetary policy is constrained, as the feedback loop between unemployment risk and aggregate demand is significantly strengthened at such times.

I show the importance of the flight-to-liquidity mechanism by comparing the results from this two-asset model to those from various models where households have access to one liquid asset. Even if such models are calibrated to match the same decline in consumption during unemployment as in the data, unemployment risk and unemployment insurance have little effect on business cycle volatility, as they lack the flight-to-liquidity and decline in investment demand that occurs in the two-asset model.

The two-asset model can also be used to study policies that affect asset adjustment costs. I consider the effect of various policies that temporarily reduce the tax faced by individuals that make a withdrawal from their illiquid asset holdings, motivated by a policy that was included in the CARES Act in response to the COVID-19 pandemic. The model implies that the particular policy included in the CARES Act, a short-lived and quickly-reversed removal of the withdrawal tax, was contractionary: such a reduction in withdrawal costs leads to a synchronized withdrawal of illiquid assets and a significant decline in investment.

A. Literature Review

The empirical evidence on the consumption response to unemployment spells in this paper builds on previous work by Gruber (1997), Aguiar and Hurst (2005), Chodorow-Reich and Karabarbounis (2016), Kolsrud et al. (2018) and others. These papers either estimate the average consumption decline during unemployment or focus on heterogeneity related only to liquid asset holdings. I provide evidence that illiquid asset holdings are also an important determinant of the response of household consumption to unemployment.

The evidence on illiquid asset withdrawal during unemployment spells is consistent with contemporaneous work by Coyne, Fadlon and Porzio (2022). While I use data from the SCF, they find a similar relationship between unemployment and penalized withdrawals from retirement accounts using U.S. tax records. While the SCF has drawbacks relative to such administrative data, offering a much smaller sample with more measurement error, it also has some benefits: I am able to use questions about asset holdings and the length of unemployment spells to show that the relationship between unemployment and early withdrawals from retire-

ment accounts is primarily driven by individuals whose unemployment spell is long or whose unemployment spell occurs when they have few liquid assets.

Turning to the evidence on precautionary saving, Basten, Fagereng and Telle (2016) use Norwegian administrative data to provide evidence of an increase in saving and a shift toward safe assets in the years leading up to an unemployment spell. While they do not study the effect of unemployment risk on illiquid asset holdings, their evidence of an increase in saving overall is consistent with my results from the SCE. To the best of my knowledge, my paper is the first to provide empirical evidence of the effect of unemployment risk on precautionary saving and portfolio allocation using a direct measure of job loss expectations.

A number of recent papers have investigated the source of the pro-cyclical skewness of the income growth distribution documented by Guvenen, Ozkan and Song (2014). I use the Current Population Survey (CPS) to show that this is driven entirely by cyclicalities in the distribution of changes in time employed. For workers that do not experience unemployment, the skewness of the income growth distribution exhibits no cyclicalities. This is consistent with evidence from Italian data provided in Hoffmann and Malacrino (2019).

This paper also contributes to the literature studying the aggregate implications of unemployment risk in HANK models. Previous papers have disagreed on whether or not unemployment risk amplifies business cycle dynamics in such models. For example, Ravn and Sterk (2017) study a one-asset model without capital. Their baseline version of the model features a degenerate wealth distribution and has the feature that unemployment risk strongly amplifies business cycle fluctuations. On the other hand, Gornemann, Kuester and Nakajima (2016) study a one-asset model with capital that is calibrated to match total household wealth. In their model, where households hold a large quantity of liquid assets, business cycle volatility is unaffected by unemployment risk: unemployment risk leads to more volatile consumption but less volatile investment.

While the above results suggest that the presence of capital is the key determinant of whether or not unemployment risk amplifies the response to aggregate shocks, it is unfortunately not that simple. Challe et al. (2017) estimate a HANK model featuring liquid capital. They achieve tractability by assuming that risk-sharing takes place between employed households, such that the wealth distribution converges to a small number of mass points. They find that unemployment risk amplifies business cycle dynamics in response to some aggregate shocks but not others.³ Finally, one other paper that studies unemployment risk in a model with multiple assets is Den Haan, Rendahl and Riegler (2017). The key difference between my model and theirs is that both assets in their model are liquid.⁴ Whether

³Cho (2023) studies an extension of their model where the liquid wealth distribution does not have finite support and finds very little amplification from unemployment risk.

⁴The two-assets in their model are bonds and equity in the firms that post vacancies in the labor

or not unemployment risk amplifies business cycles in their model depends crucially on the degree of nominal wage stickiness.

Taken together, these papers suggest that whether or not unemployment risk amplifies business cycles depends crucially on the asset structure of the economy. In this paper, I provide a quantitative assessment of the amplification provided by unemployment risk in a model that matches the distribution of both liquid and illiquid asset holdings, as well as new evidence on the consumption response to unemployment spells, the relationship between unemployment and illiquid asset adjustment, and the relationship between unemployment risk and saving in both liquid and illiquid assets. I find that unemployment risk does significantly amplify business cycle fluctuations in this environment.

By studying a HANK model with liquid and illiquid assets, my paper is closely related to the burgeoning literature following Kaplan, Moll and Violante (2018). My paper introduces search frictions in the labor market of such models. The flight-to-liquidity mechanism in the two-asset model in my paper is related to that studied by Bayer et al. (2019). In a two-asset model with a competitive labor market, they show that uncertainty shocks to households' idiosyncratic productivity can lead to a decline in investment and output. The mechanism in their model is operative in response to exogenous uncertainty shocks to idiosyncratic productivity. In my model, income risk is endogenous, as any shock that affects the unemployment rate also affects household income risk.

Finally, this paper contributes to the literature studying the role of unemployment insurance as an automatic stabilizer, such as Kekre (2016), McKay and Reis (2016) and McKay and Reis (2021). The latter uses a model with a degenerate wealth distribution, as in Ravn and Sterk (2017), and argues that the optimal level of unemployment insurance is substantially higher when taking into account its ability to dampen aggregate fluctuations. My contribution to this literature is to show that unemployment insurance can also affect business cycle volatility in a model with liquid and illiquid assets, through its ability to dampen the flight-to-liquidity that occurs when unemployment rises.

The rest of the paper is organized as follows. Section II shows that the consumption response to unemployment spells depends on both liquid and illiquid asset holdings. Section III documents the relationship between unemployment and the withdrawal of illiquid assets. Section IV provides new evidence on the effect of unemployment risk on household saving behavior. Section V describes the two-asset model, and Section VI shows that it is consistent with the empirical evidence. Section VII studies the impact of an aggregate productivity shock in the two-asset model and compares this with the results from one-asset models. Section VIII studies a temporary removal of the withdrawal tax on illiquid assets, as in the CARES Act. Section IX concludes.

market.

II. Consumption Response to Unemployment Spells

In this section, I show that the decline in household consumption during unemployment spells depends on both liquid and illiquid asset positions.

Methodology As in Kaplan, Violante and Weidner (2014), I classify households as non hand-to-mouth if they have significant liquid asset holdings, wealthy hand-to-mouth if they have few liquid assets but significant illiquid asset holdings, and poor hand-to-mouth if they have few liquid or illiquid assets. I then estimate the response of consumption to unemployment spells using the following specification:

$$(1) \quad \log C_{i,t} = \beta \mathbf{X}_{i,t} + \gamma_N U_{i,t} \mathbb{1}\{\text{N-HTM}\} + \gamma_W U_{i,t} \mathbb{1}\{\text{W-HTM}\} + \gamma_P U_{i,t} \mathbb{1}\{\text{P-HTM}\} + \epsilon_{i,t}$$

where $C_{i,t}$ denotes household consumption, $U_{i,t} \in [0, 1]$ denotes the fraction of the year that the household spent unemployed, and the indicator variables denote the liquid/illiquid asset status of the household. Chodorow-Reich and Karabarbounis (2016) use the same specification to estimate the average consumption response to unemployment spells, without conditioning on asset positions. As in their specification, I include an extensive set of controls in $\mathbf{X}_{i,t}$: region-year fixed effects; race, education, age and age squared of the household head; family size and family size squared; housing tenure; number of cars owned by the household; rental value of the household's home (split into deciles by region and year); hand-to-mouth status; and the fraction of the year spent out of the labor force.

The coefficients γ_N , γ_W , and γ_P measure the decline in log consumption during unemployment for households that are either non-hand-to-mouth, wealthy-hand-to-mouth, or poor-hand-to-mouth. Using cross-sectional variation to identify the consumption decline during unemployment relies on the assumption that the set of control variables is large enough to eliminate any omitted variable bias coming from a correlation between unemployment spells and unobservables. As an alternative, I estimate the following panel regressions based on within-household variation in consumption:

$$(2) \quad \Delta \log C_{i,t} = \alpha_t + \gamma_N \Delta U_{i,t} \mathbb{1}\{\text{N-HTM}\} + \gamma_W \Delta U_{i,t} \mathbb{1}\{\text{W-HTM}\} + \gamma_P \Delta U_{i,t} \mathbb{1}\{\text{P-HTM}\} + \Delta \epsilon_{i,t}$$

Data To estimate equation 1, I use data from the Consumer Expenditure Survey (CEX) for the period from 1996 to 2017, restricting the sample to households whose head is between the ages of 25 and 55. I follow Chodorow-Reich and Karabarbounis (2016) and use a consumption measure that covers spending on non-durables and services. The CEX measures liquid asset holdings well, but has little information on illiquid asset holdings. I therefore use homeownership as a

proxy for positive illiquid asset holdings.⁵ I define households as hand-to-mouth if they are in the bottom 50% of the liquid asset distribution at the beginning of the year that they participate in the CEX.⁶ I then define them as wealthy hand-to-mouth if they are also homeowners, and as poor hand-to-mouth if they are not.

To estimate equation 2, I use biennial data from the Panel Study of Income Dynamics (PSID) for the period from 2005 to 2017. As in the CEX, I restrict the sample to households whose head is between the ages of 25 and 55. As well as having a shorter sample than the CEX, the PSID also includes less information on consumption: the measure I use is spending on food, clothing, recreation and vacations. On the other hand, the PSID does have more accurate information on illiquid wealth holdings: I measure illiquid wealth as housing equity plus the value of retirement accounts and define households as wealthy if their illiquid wealth is greater than zero. Given the biennial nature of the PSID in this period, I group households based on their asset holdings in year $t - 2$. The Online Appendix contains further details on the construction of the datasets.

Results The results of estimating equations 1 and 2 are shown in Table 1. The estimated consumption declines are very similar using both the cross-sectional variation in the CEX and the within-household variation in the PSID. Columns 1 and 4 show estimates of the average response of consumption to unemployment without interacting unemployment with the asset indicator variables. I find that consumption is 20-25% lower during unemployment on average, in line with previous estimates.⁷

Columns 2 and 5 show the estimates when I split households only on the basis of their liquid asset holdings. The estimated consumption decline during unemployment is strongly influenced by a household's liquid asset position. Non hand-to-mouth households are able to use their liquid assets to smooth consumption during unemployment, and their consumption declines by around 15% on average. Hand-to-mouth households are less able to smooth their consumption, which declines by 25-30% on average.

Columns 3 and 6 estimate the regressions in full, now splitting hand-to-mouth households into two groups on the basis of their illiquid asset holdings. When liquid asset holdings are low, illiquid asset holdings appear to significantly affect the consumption decline during unemployment: the consumption of poor hand-

⁵In Table A.2 of the Online Appendix I use data from the Survey of Consumer Finances to show that homeownership is a good proxy for illiquid asset holdings.

⁶Kaplan, Moll and Violante (2018) report that 15% of households have negative liquid asset holdings and a further 30% of households have liquid asset holdings close to zero.

⁷A large literature has studied the average response of consumption to unemployment. Chodorow-Reich and Karabarbounis (2016) find similar estimates in both the CEX and PSID. Aguiar and Hurst (2005) use the Continuing Survey of Food Intake of Individuals (CSFII) to estimate that food expenditure falls by 19% during unemployment.

Table 1—: Consumption Response to Unemployment Spells

	CEX			PSID		
	(1)	(2)	(3)	(4)	(5)	(6)
$U_{i,t}$	-0.22 (0.015)			-0.26 (0.051)		
$U_{i,t}\mathbb{1}\{\text{N-HTM}\}$		-0.14 (0.026)	-0.14 (0.026)		-0.14 (0.069)	-0.14 (0.069)
$U_{i,t}\mathbb{1}\{\text{HTM}\}$		-0.27 (0.019)			-0.32 (0.065)	
$U_{i,t}\mathbb{1}\{\text{W-HTM}\}$			-0.23 (0.027)			-0.28 (0.131)
$U_{i,t}\mathbb{1}\{\text{P-HTM}\}$			-0.30 (0.026)			-0.34 (0.074)
$H0 : \gamma_N = \gamma_H$		0.00			0.02	
$H0 : \gamma_N = \gamma_W = \gamma_P$			0.00			0.11
$H0 : \gamma_W = \gamma_P$			0.08			0.71

Notes: Robust standard errors in parentheses. PSID standard errors are clustered by household head. Regressions weighted using sampling weights. The final three rows report the p-values for different Wald tests. CEX uses 31638 observations from 1996-2017. PSID uses 17892 observations from 2005-2017.

to-mouth households declines by at least 30%, double the decline of non hand-to-mouth households. For the wealthy hand-to-mouth, the decline is around 25%, suggesting that illiquid assets provide households with at least some ability to smooth consumption during unemployment.

To formally test the hypothesis that the size of the consumption decline depends on liquid and illiquid asset positions, Table 1 also reports the p-values of Wald tests that (1) the decline is the same for hand-to-mouth and non hand-to-mouth households, (2) the decline is the same for all three groups, and (3) that the decline for the hand-to-mouth does not depend on illiquid asset holdings. In the cross-sectional regressions using the CEX, all hypotheses can be rejected at the 10% level, confirming that both liquid and illiquid asset positions are important for determining the size of the consumption decline during unemployment. Given the smaller sample in the PSID, the second and third hypothesis cannot be rejected in the regressions using within-household variation.

One concern with the approach used here is that differences in the consumption response to unemployment spells may reflect heterogeneity in the effect of unemployment spells on household labor income, rather than heterogeneity in the effect of a given decline in labor income on household consumption. In Table A.5 of the Online Appendix I show that this is not the case: I find no evidence that the effect of a given unemployment spell on household labor income differs across

the three groups.

III. Illiquid Asset Response to Unemployment Spells

The findings in the previous section suggest that illiquid assets can play a role in smoothing consumption during unemployment spells. I now turn to data from the Survey of Consumer Finances (SCF) to understand the relationship between unemployment spells and illiquid asset holdings. I find that unemployment is a strong predictor of illiquid asset withdrawal, and that this effect is primarily driven by unemployment spells that are either long or occur when a household has few liquid assets.

Data I use data from the SCF from 2004 to 2019. To measure the withdrawal of illiquid assets, I focus on early withdrawals from tax-deferred individual retirement accounts (IRAs).⁸ Such withdrawals are generally subject to a 10% penalty, making them a clear example of illiquid asset adjustment.⁹ Along with housing equity, retirement accounts are one of the key components of illiquid asset holdings, making up around a fifth of all household wealth. I restrict the sample to households whose head is at most 55 years of age and has an IRA. More details on the sample are included in the Online Appendix.

Results Table 2 reports the annual probability of an early withdrawal for different groups of households. The first row shows that between four and five percent of households make an early withdrawal from their retirement account in a given year. The next two rows split the sample depending on whether or not the household head experienced an unemployment spell that year. Households whose head had an unemployment spell are between two and three times as likely to have made an early withdrawal from their retirement account as those whose head was employed for the whole year. This provides evidence that the withdrawal of such illiquid assets is an important way that many households smooth their consumption in the face of unemployment shocks.

Next, I further divide the sample of households whose head was unemployed into two groups, based on the length of the unemployment spell. Households whose head was unemployed for more than 12 weeks were nearly three times as likely to make an early withdrawal than those whose head was unemployed for 12 weeks or less.

Finally, I split the sample of unemployed households based on their liquid asset holdings. As in Section II, I define households as being hand-to-mouth if they are

⁸The SCF question about withdrawals from retirement accounts is specifically asked in relation to IRA/Keogh accounts, and does not relate to employer-sponsored accounts such as a 401(k).

⁹For Roth IRAs this penalty applies to earnings but not contributions. I obtain very similar results if I remove households with Roth IRAs from the sample.

Table 2—: Illiquid Asset Withdrawal Probabilities

	Data	95% C.I.	p-value
Full Sample	0.046	(0.039,0.054)	
No Unemployment Spell	0.040	(0.033,0.049)	0.000
Unemployment Spell	0.105	(0.073,0.140)	
Short Unemployment Spell	0.058	(0.026,0.098)	0.007
Long Unemployment Spell	0.159	(0.104,0.220)	
Unemployment Spell & Non-HTM	0.048	(0.011,0.099)	0.007
Unemployment Spell & HTM	0.134	(0.083,0.191)	

Notes: Probabilities constructed using sampling weights from households in the 2004 to 2019 waves of the SCF. I define an unemployment spell as short for households whose head was unemployed for 12 weeks or less. I define a household as hand-to-mouth if they have less than the median level of liquid assets. The first three sections use a sample of 4863 households. The last section uses a sample of 3649 households. Bootstrapped confidence intervals in parentheses. p-values calculated using Fisher's exact test.

in the bottom 50% of the liquid asset distribution. The last two rows of Table 2 show that households with low liquid asset holdings were almost three times as likely to make an early withdrawal than those who had high liquid asset holdings if their head had an unemployment spell. Overall, these results are consistent with the idea that liquid assets are the primary source of consumption smoothing for unemployed households, but that illiquid assets are also used when households have depleted their liquid asset holdings.

The second column of Table 2 provides bootstrapped confidence intervals for each of these probabilities, while the third column reports the p-value for tests that the probability of withdrawal does not depend on employment status, the length of the unemployment spell, or liquid asset holdings. In all cases, the null hypothesis that withdrawal probabilities are the same across the two groups can be rejected at the 1% level.

The withdrawal probabilities in Table 2 do not control for other observable variables that may be correlated with an individual's employment status and their withdrawal probability. In the Online Appendix I show that the results are unaffected by the addition of controls for age, education, race, family size and year.

IV. Precautionary Response to Unemployment Risk

The evidence in the previous two sections shows that illiquid assets provide some ability for households to smooth consumption during unemployment spells, and that households turn to these assets primarily when they have exhausted their liquid asset holdings. I now turn to considering how unemployment risk affects household behavior during employment rather than unemployment.

The findings in Sections II and III suggest that both the level of saving and the choice of which assets to save in should be closely related to expectations about the risk of becoming unemployed. To assess this, I use data from Survey of Consumer Expectations (SCE) administered by the Federal Reserve Bank of New York. Importantly for my purposes, this survey contains data on individual's perceptions of the unemployment risk that they face.

I find that consumption spending declines significantly when a household head's unemployment risk rises. This suggests that households respond to higher unemployment risk by increasing their precautionary saving. I then show that an increase in unemployment risk decreases the amount of their income that household heads plan to invest in their retirement accounts, an important form of illiquid wealth, as discussed in Section III. Thus, households both increase their saving overall, and decrease their saving in illiquid assets, in response to higher unemployment risk. Overall, this is consistent with a precautionary flight-to-liquidity motive when unemployment risk rises.

Data I use data from the SCE from 2014 to 2019. The survey occurs monthly, and interviews roughly 1,300 households heads each month. Respondents participate in the survey for up to twelve months. In addition to the main survey, which is fielded monthly, I use data from the Household Spending supplement, which is fielded every four months, as well as the Household Finance supplement, which was fielded once a year between 2014 and 2019.

From the main survey, I obtain estimates of perceived unemployment risk: employed household heads are asked every month to report their estimate for the probability that they will lose their job in the following 12 months. I denote this $\mathbb{E}_{i,t}[s_{t,t+12}^i]$. I am interested in understanding how household consumption and saving decisions are affected by unemployment risk, both in terms of the level of saving overall and also in terms of how unemployment risk affects household portfolio decisions. Consequently, from the Household Spending supplement I obtain data on household consumption. Respondents are first asked whether their current consumption spending is higher or lower than it was 12 months ago. They are then asked a second question to put the change in consumption spending in percentage terms. I use an indicator, $\mathbb{1}\{c_t^i \geq c_{t-12}^i\}$, to denote that individual i reports consumption spending that is "about the same or higher" than 12 months ago. I denote the percentage change in consumption from 12 months ago by $\frac{c_t^i - c_{t-12}^i}{c_{t-12}^i} \cdot 10$.

Finally, from the Household Finance supplement, I obtain data on household asset allocation. Specifically, respondents who report having Defined Contribution

¹⁰A very small number of individuals report extreme values when asked to put a percentage value on their annual change in consumption. Thus, I winsorize the 0.5% largest values of this variable.

Table 3—: Precautionary Response to Unemployment Risk

	(1) $\mathbb{1}\{c_t^i \geq c_{t-12}^i\}$	(2) $\frac{c_t^i - c_{t-12}^i}{c_{t-12}^i}$	(3) $\mathbb{1}\{\mathbb{E}_{i,t}[P_{t,t+12}^i > P_{t-12,t}^i]\}$
$\Delta_9 \mathbb{E}_{i,t}[s_{t,t+12}^i]$	-.28 (.073)	-.038 (.014)	-.32 (.13)
Observations	946	942	260

Notes: Estimates from estimating equation 3. Dependent variable shown in the second row. c_t^i denotes household consumption of individual i in month t . $\mathbb{1}\{\mathbb{E}_{i,t}[P_{t,t+12}^i > P_{t-12,t}^i]\}$ is an indicator denoting that individual i expects to increase the proportion of earnings contributed to their Defined Contribution pension over the next year. $\Delta_9 \mathbb{E}_{i,t}[s_{t,t+12}^i]$ denotes the change in the perceived annual job loss probability of individual i from month $t-9$ to t . Robust standard errors in parentheses. Regressions weighted using sampling weights. In column (2) I winsorize to remove the most extreme 0.5% of consumption growth observations. Data from the Survey of Consumer Expectations from 2014 to 2019.

pension accounts are asked whether they expect to increase, decrease, or leave unchanged the proportion of their earnings that they contribute to this pension account in the next 12 months, relative to the last 12 months. I summarize this by an indicator, $\mathbb{1}\{\mathbb{E}_{i,t}[P_{t,t+12}^i > P_{t-12,t}^i]\}$, which denotes that individual i expects to increase their pension contributions in the next year.

Using this data, I run regressions of the following form:

$$(3) \quad Y_{i,t} = \beta \mathbf{X}_{i,t} + \gamma \Delta_9 \mathbb{E}_{i,t}[s_{t,t+12}^i] + \epsilon_{i,t}$$

$Y_{i,t}$ is the outcome of interest: one of the measures of consumption changes or the expected change in pension contributions. $\Delta_9 \mathbb{E}_{i,t}[s_{t,t+12}^i]$ is the change in the perceived job loss probability over the past 9 months, and $\mathbf{X}_{i,t}$ is a vector of control variables: time fixed effects; age, age squared, gender and education of the household head; and the household's SCE income group. I use the 9-month change in job loss expectations in order to investigate the effect of a persistent change in perceived unemployment risk while preserving the size of the sample.

As in the SCF, I restrict the sample to households whose head is at most 55 years of age. I also only include individuals who report having been employed in the same job for at least one year. This bolsters my confidence that the results on changes in consumption expenditure are due to changes in precautionary saving, rather than changes in household income. More details are included in the Online Appendix, where I also show the robustness of my results to using the 10- or 11-month change instead.¹¹

¹¹Using the 11-month change shrinks the sample by roughly two-thirds relative to the 9-month change. The small size of the SCE sample is particularly notable for the regression using the Household Finance supplement, given that this was only administered annually for six years.

Results Table 3 provides the results of estimating equation 3 with three different dependent variables. First, in column (1), the dependent variable is the indicator for household consumption spending being higher than a year before. I find that an increase in perceived unemployment risk significantly lowers the probability that household consumption spending has increased. The standard deviation of the 9-month change in perceived unemployment risk is 20%. Thus, a one standard deviation increase in perceived unemployment risk is associated with a roughly 6 percentage point lower likelihood that household spending has increased.¹²

In column (2) the dependent variable is the reported percentage change in consumption spending. Using this measure, I am able to quantify the decline in spending when unemployment risk rises. The estimated coefficient implies that a 20% rise in perceived unemployment risk is associated with a 0.7-0.8% decline in consumption spending.

Thus, the results in column (1) and column (2) provide evidence that households cut back their consumption spending, and thus increase their saving, when unemployment risk rises. To understand whether households increase their liquid or illiquid asset holdings, in column (3) the dependent variable is an indicator that is equal to one if the individual expects to invest a higher proportion of their earnings in their pension in the next year, relative to the previous year. The estimated coefficient implies that a 20% rise in perceived unemployment risk leads to a 6 percentage point decline in the probability that the individual expects to increase their pension contributions. Roughly 30% of individuals report that they expect to increase the proportion of their earnings that they contribute on average, so this suggests that higher perceived unemployment risk leads to a quantitatively significant decline in illiquid asset investment.


In conclusion, evidence from the SCE shows that households respond to increased unemployment risk by saving more overall, but by saving less in illiquid forms of wealth. This is consistent with the evidence from the previous sections, which suggested that liquid assets are preferable for consumption smoothing during unemployment spells.

V. A Two-Asset Model with Unemployment

Motivated by the empirical evidence in the previous sections, I now study the role of endogenous unemployment risk in a heterogeneous-agent New Keynesian model with both liquid and illiquid assets. As in Kaplan, Moll and Violante (2018), households trade both liquid assets (nominal bonds) and illiquid assets

¹²On average, individuals report a 14% chance that they will lose their job in the next year. 80% of individuals report consumption being “about the same or higher” than a year before.

Figure 1. : Model Timeline



(physical capital).¹³ Search frictions in the labor market render unemployment, and consequently income risk, endogenous to aggregate shocks.

In the model, households face a trade-off when choosing their asset portfolio. Bond holdings can be adjusted without cost, but offer a low rate of return. Capital offers a higher return, but is costly to adjust. As bonds are liquid, they are well suited to smoothing consumption in response to transitory income shocks, such as unemployment spells. The key mechanism in this model is that a household's optimal asset portfolio depends on the level of unemployment risk in the economy, leading to a time-varying preference for holding liquid assets.

Households Time is discrete. There is a continuum of infinitely-lived households that supply labor inelastically, derive utility from consumption, and trade both liquid and illiquid assets. Households' idiosyncratic labor productivity follows an exogenous Markov process. Households are also subject to shocks to their employment status. In each period, households that choose to adjust their illiquid asset position pay a random adjustment cost, described in more detail below. In order to study the effect of the CARES Act, which reduced the tax on early withdrawals from retirement accounts, I also assume that households that withdraw from their illiquid asset holdings are required to pay a tax, τ_k . At the end of each period, a fraction ζ of households dies and is replaced by new households with zero wealth.¹⁴ The government receives the assets of the dying households, effectively levying a 100% estate tax. Within a period, the timing of events is shown in Figure 1.

For households that choose to adjust their illiquid asset holdings, the recursive

¹³I follow Kaplan, Moll and Violante (2018) in calibrating the model assuming that housing, business equity, and durables are illiquid assets. In an earlier version of this paper, Graves (2020), I showed that the main results are the same if the illiquid asset in the model is housing rather than productive capital.

¹⁴As in Kaplan, Moll and Violante (2018), this assumption helps the model to match the wealth distribution in the data.

problem is:

$$(4) \quad V_t^A(b, k, z, e) = \max_{c, b', k'} \frac{c^{1-\gamma}}{1-\gamma} + \beta(1-\zeta)\mathbb{E}_{e', z'} V_{t+1}(b', k', z', e')$$

subject to

$$(5) \quad \begin{aligned} k' + b' + c + \tau_k \mathbb{1}\{k' < k\}(k - k') &= \mathbb{1}\{e = E\}w_t z(1 - \pi) + \mathbb{1}\{e = U\}w_t \phi(z)(1 - \pi) \\ &\quad + R_t^b(b)b + R_t^k k + T_t \\ b' &\geq -\underline{b} \\ k' &\geq 0 \\ z' &= \Gamma(z) \end{aligned}$$

where b and k denote bond and capital holdings, z is the household's idiosyncratic productivity, and $e \in \{E, U, N\}$ is the household's employment status, equal to E if employed, U if unemployed and receiving unemployment insurance, or N if unemployed and not receiving unemployment insurance. If employed, the household receives wage w_t per unit of labor productivity. If unemployed and receiving unemployment insurance, $e = U$, households receive benefits equal to $w_t \phi(z)$ per unit of labor productivity. Both sources of labor income are subject to a linear tax, τ_l . If unemployed and not receiving unemployment insurance, $e = N$, the household received no labor income. However, T_t denotes a lump-sum transfer which is received by all households.

Households face borrowing constraints on their holdings of both liquid bonds and illiquid capital. Illiquid asset holdings must be non-negative. Household are able to borrow up to \underline{b} units of the liquid asset. However, there is an exogenous wedge, κ , between the borrowing and lending rates on the liquid asset:¹⁵

$$(6) \quad R^b(b) = \begin{cases} \frac{1+i_t}{\Pi_t} & \text{if } b \geq 0 \\ \frac{1+i_t}{\Pi_t} + \kappa & \text{if } b < 0 \end{cases}$$

where i_t is the nominal interest rate set by the central bank, and Π_t is the gross rate of inflation. The return on the illiquid asset is derived from supplying capital services to the intermediate good producers at rate r_t^k . Capital services provided are the product of the utilization rate, u_t , and the household's holding of the illiquid asset, k . The rate of depreciation of capital is increasing in the utilization rate, as in Greenwood, Hercowitz and Huffman (1988).¹⁶ Thus, the rate of return

¹⁵This assumption helps to ensure a realistic distribution of liquid asset holdings: a large mass of households with close to zero liquid assets, and a share of around 15% of households with negative liquid asset holdings.

¹⁶Without variable capital utilization, the marginal product of labor, and thus the wage, would rise after a negative productivity shock that leads to a decline in employment and a consequent rise in the

on the illiquid asset is:

$$(7) \quad R_t^k = 1 + r_t^k u_t - \delta_0 u_t^{\delta_1}$$

If the household doesn't adjust their illiquid asset holdings, their problem is:

$$(8) \quad V_t^{NA}(b, k, z, e) = \max_{c, b'} \frac{c^{1-\gamma}}{1-\gamma} + \beta(1-\zeta)\mathbb{E}_{z', e'} V_{t+1}(b', k, z', e')$$

subject to

$$k + b' + c = \mathbb{1}\{e = E\}w_t z(1 - \tau_l) + \mathbb{1}\{e = U\}w_t \phi(z)(1 - \tau_l) + R_t^b(b)b + R_t^k k + T_t$$

$$b' \geq -\underline{b}$$

$$z' = \Gamma(z)$$

Illiquid Asset Adjustment Costs Each period, household's draw an iid adjustment cost, χ , from the uniform distribution on $[0, \bar{\chi}]$, denominated in units of utility. They then decide whether or not to adjust their capital holdings. Consequently, the value of the household's problem, conditional on a draw of χ is:

$$(9) \quad V_t(b, k, z, e; \chi) = \max\{V_t^A(b, k, z, e) - \chi, V_t^{NA}(b, k, z, e)\}$$

The value before the draw of χ is:

$$(10) \quad V_t(b, k, z, e) = \mathbb{E}_\chi V_t(b, k, z, e; \chi)$$

Such random adjustment costs have been used in the household context by Bayer et al. (2019) and in the firm context by Khan and Thomas (2008). When calibrating the model, I discipline the size of these adjustment costs using data on liquid and illiquid wealth holdings.

Idiosyncratic Shocks Households face idiosyncratic shocks to their employment status and to their productivity. Each period, employed households are separated to unemployment with exogenous probability s . Unemployed households find employment with endogenous probability f_t . If unemployed, the probability that households receive unemployment insurance is independent across periods and equal to ξ .¹⁷ I assume that households whose employment is terminated may

capital-labor ratio.

¹⁷I assume that reciprocity is random as there is no evidence in the SCF that reciprocity is related

immediately re-enter employment.

Previous research has shown that having a realistic income process is crucial if models are to generate a realistic wealth distribution. A key feature of the data is the high level of kurtosis of the income growth distribution. By introducing infrequent large income changes, idiosyncratic unemployment risk helps to provide high kurtosis of income growth. However, to match the level seen in the data, I also assume that idiosyncratic productivity shocks are infrequent. Specifically:

$$(11) \quad \log z' = (1 - \rho_z)\mu_z + \rho_z \log z + \epsilon_z$$

$$(12) \quad \epsilon_z = \begin{cases} N(0, \sigma_z^2) & \text{with prob } \lambda_z \\ 0 & \text{with prob } 1 - \lambda_z \end{cases}$$

I introduce the normalization μ_z to ensure that the mean value of idiosyncratic productivity is equal to 1.

Final Good Producers There is a representative final good producer, which aggregates a continuum of intermediate goods according to the production function:

$$(13) \quad Y_t = \left(\int_0^1 y_{j,t}^{\frac{\epsilon-1}{\epsilon}} dj \right)^{\frac{\epsilon}{\epsilon-1}}$$

Their profit maximization problem leads to the following demand curve for intermediate goods:

$$(14) \quad y_{j,t}(p_{j,t}) = \left(\frac{p_{j,t}}{P_t} \right)^{-\epsilon} Y_t$$

$$(15) \quad P_t = \left(\int_0^1 p_{j,t}^{1-\epsilon} dj \right)^{\frac{1}{1-\epsilon}}$$

Intermediate Good Producers Intermediate goods are produced using both capital services, $k_{j,t}$, and labor, $n_{j,t}$, using the production function:

$$(16) \quad y_{j,t} = A_t k_{j,t}^\alpha n_{j,t}^{1-\alpha}$$

where A_t is the level of aggregate productivity. Intermediate good producers rent

to liquid asset holdings: 48% of households that have unemployment spells report receiving unemployment insurance, while the proportion is 50% for households with low liquid asset holdings and 45% for households with high liquid asset holdings.

capital from households at rate r_t^k and labor from a representative labor agency at rate h_t . Their cost minimization problem implies the following value for their marginal cost of production:

$$(17) \quad m_t = \frac{1}{A_t} \left(\frac{r_t^k}{\alpha} \right)^\alpha \left(\frac{h_t}{1-\alpha} \right)^{1-\alpha}$$

I assume that intermediate good producers are owned by risk-neutral entrepreneurs who consume all profits each period. In the Online Appendix I show that the results are similar if instead profits are distributed lump-sum to households. Price adjustment is subject to quadratic costs.¹⁸ Given these assumptions, the recursive form of their price-setting problem is:

$$(18) \quad V_t^I(p_{j,t-1}) = \max_{p_{j,t}} Y_t \left\{ \left(\frac{p_{j,t}}{P_t} - m_t \right) \left(\frac{p_{j,t}}{P_t} \right)^{-\epsilon} - \frac{\theta_P}{2} \log \left(\frac{p_{j,t}}{p_{j,t-1}} \right)^2 \right\} + \beta V_{t+1}^I(p_{j,t})$$

where θ_P governs the size of price adjustment costs. The solution to this problem implies the following New Keynesian Phillips Curve:

$$(19) \quad \log(\Pi_t) = \beta \frac{Y_{t+1}}{Y_t} \log(\Pi_{t+1}) + \frac{\epsilon}{\theta_P} (m_t - m^*)$$

where $m^* = \frac{\epsilon-1}{\epsilon}$ is the inverse of the steady-state mark-up and $\Pi_t = \frac{P_t}{P_{t-1}}$.

Labor Agency Intermediate good producers hire labor from a representative labor agency. This agency hires households in a frictional labor market by posting vacancies. I assume that the labor agency is also owned by risk-neutral entrepreneurs. The labor agency's recursive problem is:

$$(20) \quad J_t(N) = \max_{N', V} (h_t - w_t)N' - cV + \beta J_{t+1}(N')$$

subject to

$$N' = (1-s)N + q(\theta_t)V$$

where N is the number of employed households, V is the number of vacancies, c is the cost of posting a vacancy, $q(\theta_t)$ is the job-filling probability, and $\theta_t \equiv \frac{V_t}{U_t}$ is labor market tightness.

There are two wages in the model: h_t is the wage paid by intermediate good

¹⁸As in Rotemberg (1982).

producers to the labor agency, and w_t is the wage paid by the labor agency to employed households. Due to the search frictions in the model, there is a range of household wages that is between the reservation wages of households and the labor agency. I assume the following wage rule, which implies that the wage paid to households responds to the wage paid to the labor agency with elasticity ϵ_w .¹⁹

$$(21) \quad w_t = \bar{w} \left(\frac{h_t}{\bar{h}} \right)^{\epsilon_w}$$

Labor Market The labor market is characterized by search and matching frictions. Given U_t unemployed households and V_t vacancies, $M(U_t, V_t)$ new employment relationships are formed according to the following matching function:²⁰

$$(22) \quad M(U_t, V_t) = \frac{U_t V_t}{(U_t^l + V_t^l)^{\frac{1}{l}}}$$

The job-finding and job-filling rates are functions of labor market tightness:

$$(23) \quad f(\theta_t) = (1 + \theta_t^{-l})^{-\frac{1}{l}}$$

$$(24) \quad q(\theta_t) = (1 + \theta_t^l)^{-\frac{1}{l}}$$

Fiscal and Monetary Policy The central bank sets nominal interest rates according to the following Taylor rule:

$$(25) \quad i_{t+1} = \bar{r}^b + \psi \log(\Pi_t)$$

Unemployment insurance provides a replacement rate ϕ_0 and is capped at a fraction ϕ_1 of the average wage:

$$(26) \quad \phi(z) = \min\{\phi_0 z, \phi_1\}$$

The government receives revenue from the labor income tax, the illiquid asset withdrawal tax, and the estate tax. It distributes unemployment insurance and the lump-sum transfer, issues nominal bonds, and undertakes government spend-

¹⁹The complexity of the problem precludes a Nash bargaining solution for wages. Similar wage rules are used in Gornemann, Kuuster and Nakajima (2016) and Den Haan, Rendahl and Riegler (2017). In the Online Appendix I show that the main results of the paper are robust to a wide range of values of ϵ_w .

²⁰As in Den Haan, Ramey and Watson (2000). This matching function ensures that job-finding and job-filling rates are well-defined for any value of $\theta_t > 0$.

ing. The government budget constraint is:

(27)

$$\begin{aligned} G_t + r_t^b B_t^g + T_t + \xi(1 - N_t)w_t \int \phi(z) d\mu_t &= \tau_l N_t w_t + \tau_l \xi(1 - N_t)w_t \int \phi(z) d\mu_t \\ &+ \tau_k \int \mathbf{1}\{k' < k\}(k - k') d\mu_t \\ &+ \zeta \int (R_t^b(b)b + R_t^k k) d\mu_t \end{aligned}$$

Equilibrium An equilibrium in this model consists of paths for household decision rules $\{c_t, b_t, k_t, u_t\}_{t=0}^\infty$, firm decision rules $\{L_t, K_t, N_t, V_t\}_{t=0}^\infty$, prices and returns $\{w_t, h_t, r_t^b, r_t^k\}_{t=0}^\infty$, inflation $\{\Pi_t\}_{t=0}^\infty$, the job finding rate $\{f_t\}_{t=0}^\infty$, fiscal variables $\{G_t, T_t, B_t\}_{t=0}^\infty$, and the distribution of households $\{\mu_t\}_{t=0}^\infty$ such that:

- 1) Decision rules solve household and firm problems, taking as given aggregate variables
- 2) The government budget constraint holds
- 3) The distribution satisfies aggregate consistency conditions
- 4) All markets clear

Market Clearing The following market clearing conditions must hold in equilibrium:

- 1) Bonds:

$$(28) \quad B_t^g = B_t^h = \int b d\mu_t$$

- 2) Capital:

$$(29) \quad K_t = K_t^h = u_t \int k d\mu_t$$

- 3) Labor:

$$(30) \quad L_t = N_t = \int \mathbf{1}\{e = E\} d\mu_t$$

4) Goods:

$$(31) \quad Y_t = C_t + I_t + G_t + \Theta_t + \kappa \int \max\{-b, 0\} d\mu_t + cV_t$$

The goods market clearing condition takes into account price adjustment costs, Θ_t , as well as the borrowing costs and costs of posting vacancies.

A. Calibration

Table 4 summarizes the calibration of the model. The model period is one quarter. Below, I provide further details on the calibration process.

Labor Market The quarterly job separation rate is 0.1, in line with the Job Openings and Labor Turnover Survey (JOLTS). I target a steady-state unemployment rate of 6%, and a quarterly job-filling rate of 0.71, as in Den Haan, Ramey and Watson (2000). These values imply a matching function elasticity of $l = 1.68$. I set the vacancy cost equal to 5% of the quarterly wage. Combined with the job-filling probability, this implies a hiring cost per worker of around 7% of the quarterly wage, as in Christiano, Eichenbaum and Trabandt (2016). With this assumption, I calibrate the steady-state wage to generate an unemployment rate of 6%. I set ϵ_w to 0.45, the elasticity of wages to labor productivity estimated by Hagedorn and Manovskii (2008).²¹

Income Process I set the values of ρ_Z , σ_Z , and λ_Z in order to target the variance and kurtosis of the annual income growth distribution, as well as the variance of the level of income. The top section of Table 5 reports these moments in the model and the data. While the high kurtosis of the income growth distribution implies that idiosyncratic productivity shocks occur infrequently, unemployment spells provide income shocks that are both more frequent and more transitory.

Wealth Distribution The key parameters affecting the liquid and illiquid wealth distributions are the coefficient of relative risk aversion, the death rate, the discount factor, the borrowing wedge, and the parameter governing the degree of illiquid asset adjustment costs. I set the coefficient of relative risk aversion, γ , to 2 and the quarterly death probability, ζ , to $\frac{1}{180}$, as in Kaplan, Moll and Violante (2018), implying that households live for 45 years on average. I calibrate the other parameters to target the total quantity of liquid and illiquid assets

²¹Due to movements in the mark-up, this calibration leads to wages that are more responsive to labor productivity than in the data. This ensures that the results of the model are not driven by the stickiness of real wages, as further shown in the Online Appendix.

relative to output, as well as the fraction of households with negative liquid asset holdings, as reported by Kaplan, Moll and Violante (2018). Table 5 provides various moments of the wealth distribution.²²

The model matches the Gini coefficient for total wealth inequality. The model is also close to matching the proportion of hand-to-mouth households, defined as those with liquid asset holdings close to zero. The bottom two panels of Table 5 provide further details on the share of the liquid and illiquid wealth distributions held by different quantiles. The model slightly fails to match the wealth holdings of the top 1% of households, and instead over-predicts the share of wealth held by the rest of the top 10% of the distribution. In terms of adjustment probabilities, 3% of employed households and 11% of unemployed households adjust their illiquid asset holdings each period. The total adjustment costs that households pay are equivalent to 0.7% of aggregate output.²³

Fiscal and Monetary Policy The particular details of unemployment insurance vary across US states. I set the cap on unemployment insurance, ϕ_1 , to two-thirds of the average wage, and the replacement rate, ϕ_0 , to 50%. These values are the most common across states, as reported in Department of Labor (2018). The parameter ξ governs the probability that unemployed households receive unemployment insurance. Figure A.1 in the Online Appendix shows that a large fraction of unemployed individuals do not actually receive unemployment insurance, even if their unemployment spell is short enough to qualify for benefits. I set ξ equal to 0.45, the average UI reciprocity rate for the short-term unemployed. I set the illiquid asset withdrawal tax to 10%, equal to the standard penalty for early withdrawals from retirement accounts. I set the linear income tax to 30%, and the value of the transfer to 0.04, such that it is equal to around 1% of GDP, as in McKay and Reis (2016). I set the steady-state real return on bonds to 1% on an annual basis. I assume that the Taylor rule coefficient on inflation is 1.5.

Remaining Parameters I calibrate the remaining parameters of the model to standard values in the New Keynesian literature. The coefficient on capital in the intermediate good production function is set to 0.33. I choose δ_0 such that the depreciation rate on capital is 6% at an annual frequency and δ_1 such that the steady-state utilization rate is equal to 1. I set the elasticity of substitution, ϵ , to 20, implying a steady-state mark-up of 5%. I choose a low mark-up to ensure that profits are small, given that I assume that all profits are consumed by risk-

²²The target for liquid assets to output is notably lower than the ratio of government debt to GDP in the data. However, it is close to the level of government debt held by private domestic agents, which is the appropriate comparison for the model.

²³See the Online Appendix for the derivation of this value. Kaplan, Moll and Violante (2018) report that illiquid asset adjustment costs in their model total less than 4% of GDP.

Table 4—: Parameter Values

Parameter		Value	Source/Target
Separation Rate	s	0.1	JOLTS
Vacancy Cost	c	0.11	5% of Quarterly Wage
Steady-state Wage	\bar{w}	2.1	6% Unemployment Rate
Wage Elasticity	ϵ_w	0.45	Hagedorn and Manovskii (2008)
Matching Function Elasticity	l	1.68	Quarterly Job-Filling Probability
Prod. Persistence	ρ_z	0.964	Variance of Annual Income
Prod. Variance	σ_z	3.2	Variance of Annual Income Growth
Prod. Shock Probability	λ_z	0.007	Kurtosis of Annual Income Growth
Risk Aversion	γ	2	Standard value
Discount Factor	β	0.982	Illiquid Assets/Output
Death Probability	ζ	$\frac{1}{180}$	Kaplan, Moll and Violante (2018)
Adjustment Cost Limit	$\bar{\chi}$	1	Liquid Assets/Output
Borrowing Limit	\underline{b}	1	50% of Average Quarterly Labor Income
Borrowing Wedge	κ	0.019	% Negative Liquid Assets
UI Replacement Rate	ϕ_0	0.5	Department of Labor (2018)
UI Cap	ϕ_1	0.67	Department of Labor (2018)
UI Probability	ξ	0.45	Employment & Training Administration
Income Tax	τ_l	0.3	Kaplan, Moll and Violante (2018)
Withdrawal Tax	τ_k	0.1	IRA Withdrawal Penalty
Transfer	T	0.04	McKay and Reis (2016)
Return on Liquid Assets	\bar{r}^b	0.0025	1% Annual Rate of Return
Taylor Rule Coefficient	ψ	1.5	Kaplan, Moll and Violante (2018)
Capital Share	α	0.33	Standard value
Steady-State Depreciation Rate	δ_0	0.014	6% Annual Rate of Depreciation
Depreciation Elasticity	δ_1	2.02	Steady-State Utilization Rate of 1
Elasticity of Substitution	ϵ	20	Mark-up of 5%
Price Adjustment Cost	θ_P	250	Slope of New Keynesian Phillips Curve

Notes: This Table reports the baseline calibration of the two-asset model with unemployment insurance. The calibration is described in Section V.A.

Table 5—: Income and Wealth Distributions

Moment	Data	Model
Income		
Variance: Annual Log Earnings	0.70	0.71
Variance: 1-year change	0.23	0.23
Kurtosis: 1-year change	17.8	18.4
Wealth		
Liquid Assets to Output	0.26	0.30
Illiquid Assets to Output	2.92	2.88
% Poor Hand-to-Mouth	0.10	0.05
% Wealthy Hand-to-Mouth	0.20	0.20
% Negative Liquid Assets	0.15	0.14
Gini Coefficient (Total Wealth)	0.81	0.82
Top 1% share (Liquid)	47	34
Top 1%-10% share (Liquid)	39	52
Top 10%-50% share (Liquid)	18	15
Bottom 50% share (Liquid)	-4	-2
Top 1% share (Illiquid)	33	19
Top 1%-10% share (Illiquid)	37	55
Top 10%-50% share (Illiquid)	27	24
Bottom 50% share (Illiquid)	3	2

Notes: Income moments are based on Social Security Administration data on male earnings, reported by Guvenen et al. (2015). Wealth moments are from the 2004 SCF, reported by Kaplan, Moll and Violante (2018). Moments from the model are calculated by simulating 1 million households in the steady-state of the model and aggregating income to an annual frequency. In the model, I define household as hand-to-mouth if the absolute value of their liquid asset holdings is less than 10% of the average quarterly wage. I define households as wealthy if their illiquid asset holdings exceed 60% of average quarterly labor earnings.

Table 6—: Consumption Response to Unemployment Spells

	Data (CEX)			Two-Asset Model		
	(1)	(2)	(3)	(4)	(5)	(6)
$U_{i,t}$	-0.22 (0.015)			-0.18		
$U_{i,t}\mathbf{1}\{\text{N-HTM}\}$		-0.14 (0.026)	-0.14 (0.026)		-0.08	-0.08
$U_{i,t}\mathbf{1}\{\text{HTM}\}$		-0.27 (0.019)			-0.26	
$U_{i,t}\mathbf{1}\{\text{W-HTM}\}$			-0.23 (0.027)			-0.23
$U_{i,t}\mathbf{1}\{\text{P-HTM}\}$			-0.30 (0.026)			-0.27

Notes: Robust standard errors in parentheses. Regressions weighted using CEX sampling weights, with 31638 observations from 1996 to 2017. In both the model and the data, households are defined as hand-to-mouth if their liquid asset holdings are below the median. In the CEX I define households as wealthy if they are homeowners. In the model, I define households as wealthy if their illiquid asset holdings exceed 60% of average quarterly labor earnings.

neutral entrepreneurs. I then set the value of the price-adjustment cost, θ_P , to 250, which implies that the slope of the New-Keynesian Phillips curve is 0.08. If price-adjustment was of the Calvo form, this would be equivalent to prices lasting four quarters on average.

VI. Model Validation


Before turning to the effect of aggregate shocks in the model, I start by checking that the model is consistent with the empirical findings in Sections II and III. To do this, I simulate a large panel of households in the steady-state of the model and aggregate to an annual frequency. Using this panel, I run the same consumption regressions as in Section II, and calculate illiquid asset withdrawal probabilities as in Section III. I then consider the effect of idiosyncratic shocks to expected job separation probabilities, as in Section 4.

A. Consumption Response to Unemployment Spells

Table 6 compares the regression results in the model and the data. Column (4) shows that the average consumption decline during unemployment in the model is close to, but slightly smaller than, that estimated in the data. Columns (5) and (6) show that the model matches well the ranking seen in the data, including the important role for illiquid asset holdings within the hand-to-mouth group.

To understand how the two-asset model is able to generate these patterns, Figure 2 plots the log difference between the consumption of unemployed and employed

Figure 2. : Consumption Declines and Illiquid Asset Adjustment



Notes: This figure plots the log difference in consumption between unemployed and employed households at the median level of idiosyncratic productivity and mean level of illiquid asset holdings. The mean (median) value of k in the steady-state of the model is 34 (6). The mean (median) value of b is 3.4 (0.5).

households across the liquid wealth distribution. I hold illiquid wealth constant at its mean level, so the figure shows the consumption decline during unemployment for wealthy hand-to-mouth or non hand-to-mouth households.

For households that do not adjust their illiquid asset holdings, the figure shows that the consumption decline during unemployment varies strongly with a household's liquid asset position. Thus, if their liquid asset holdings are low, the consumption decline for wealthy hand-to-mouth households can be large.

On the other hand, if wealthy hand-to-mouth households do adjust their illiquid asset holdings, then their consumption decline during unemployment is negligible. In this case, the consumption decline during unemployment is largely independent of a household's liquid wealth.

Overall, the model is able to generate a realistic consumption decline for wealthy hand-to-mouth households because only a small fraction of wealthy hand-to-mouth households choose to liquidate capital during unemployment. This implies that the average consumption decline for the wealthy hand-to-mouth is between the large decline of the poor hand-to-mouth and the small decline of the non hand-to-mouth, as it in the data.

Table 7—: Illiquid Asset Withdrawal Probabilities

	SCF Data	Model
Full Sample	0.046	0.079
No Unemployment Spell	0.040	0.058
Unemployment Spell	0.105	0.185
Short Unemployment Spell	0.058	0.115
Long Unemployment Spell	0.159	0.325
Unemployment Spell & Non-HTM	0.048	0.071
Unemployment Spell & HTM	0.134	0.291

Notes: Probabilities constructed using sampling weights from the 2004 to 2019 waves of the SCF. The first three sections use a sample of 4863 households. The last section uses a sample of 3649 households. I define an unemployment spell as short for households whose head was unemployed for 12 weeks or less. I define a household as hand-to-mouth if they have less than the median level of liquid assets.

B. Illiquid Asset Response to Unemployment Spells

Table 7 compares the illiquid asset withdrawal probabilities in the model and the data. As individual retirement accounts are only one type of illiquid asset, there is no direct comparability between the levels of the withdrawal probabilities in the model and the data.²⁴ The true withdrawal probabilities in the data are higher when including withdrawals from other illiquid assets, such as housing.²⁵ However, it is possible to validate the model by considering the relative effect of unemployment and liquid asset holdings on withdrawal probabilities.


The model matches the patterns seen in the data. In the model, the withdrawal probability for households who experienced an unemployment spell is significantly higher than that for households who did not. As in the data, this is driven particularly by households who experienced long unemployment spells and those who became unemployed when they had few liquid assets.

Figure 3 plot illiquid asset withdrawal probabilities across the liquid wealth distribution to highlight the relationship between employment status and illiquid asset adjustment in the model. The withdrawal probability for employed households has a modest negative relationship with liquid asset holdings: Employed households would like to make a withdrawal from their illiquid asset holdings when their liquid wealth is low, in order to balance their portfolios and improve their ability to smooth consumption during unemployment. However, they only do so if they receive a favorable draw of the illiquid asset adjustment cost.

²⁴Also, household decisions regarding retirement accounts are intimately tied up with life-cycle considerations, from which the model abstracts.

²⁵For example, Bhutta and Keys (2016) find that an average of 11% of households extracted equity from their home each year between 1999 and 2010.

Figure 3. : Adjustment Probabilities and Employment Status



Notes: This figure plots illiquid asset withdrawal probabilities for employed and unemployed households at the median level of idiosyncratic productivity and mean level of illiquid asset holdings. The mean (median) value of k in the steady-state of the model is 34 (6). The mean (median) value of b is 3.4 (0.5).

Relative to employed households, unemployed households are much more likely to withdraw from their illiquid asset holdings when their liquid asset holdings are low. This is consistent with Figure 2 which show that such households would experience a large drop in consumption if they did not draw down upon their illiquid wealth.

C. Precautionary Response to Unemployment Risk

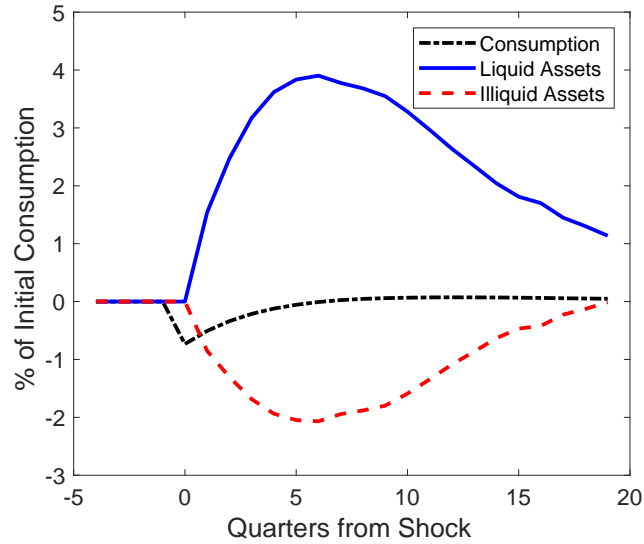
Finally, I consider the effect of idiosyncratic shocks to expected job separation probabilities. As the model does not include shocks to the job separation rate, I consider the effect of a one-time unexpected shock to the job separation rate such that the expected job loss probability over the next year rises by 20%, which I found to be the standard deviation of the change in the expected job separation rate in the SCE data in Section IV.²⁶

Figure 4 plots the average response of consumption as well as liquid and illiquid asset holdings to such a shock, relative to the paths that would have occurred with no shock to the expected job loss probability.²⁷ As in the data, I consider a panel

²⁶I do this by raising the quarterly separation rate from 0.1 to 0.21, and letting it decline back to its steady-state value as an AR(1) process with a quarterly persistence of 0.8.

²⁷I consider the effect of idiosyncratic shocks by holding all prices fixed, consistent with the use of time fixed effects in the regressions in Section IV.

Figure 4. : A Shock to the Expected Job Separation Rate



Notes: This figure plots the response of consumption and liquid/illiquid asset holdings to an unexpected shock to the job separation rate, such that the expected job loss probability over the next year rises by 20% on impact. Responses are plotted for continuously employed individuals relative to a path with no such shock. See the text for full details.

of individuals that are constantly employed. There are two main results. First, I find that consumption declines on impact by around 0.75% in response to a shock that raises the expected job loss probability over the next year by 20%. This is almost exactly in line with the magnitude of the consumption response found in the first column of Table 3. Second, I find that illiquid asset holdings fall, despite the increase in precautionary saving implied by the decline in consumption. As the model has no notion of regular contributions to a retirement account, it is not possible to quantitatively compare the decline in illiquid asset investment in the model and the data. Thus, the model is consistent both with the quantitative magnitude of the decline in consumption in response to a rise in unemployment risk, and with the fact that this increase in precautionary saving is accompanied by a shift in portfolios towards liquid wealth. In the model, both the decline in consumption and the decline in illiquid asset holdings contribute roughly equally to a sharp rise in liquid wealth.

VII. Response of the Economy to Aggregate Shocks

Having shown that the two-asset model is consistent with the new empirical evidence, I now study the response of the economy to an unanticipated negative

shock to aggregate productivity.²⁸ To understand whether or not unemployment risk affects business cycle dynamics, and if unemployment insurance is an important automatic stabilizer, I compare the impulse responses of three different versions of the model: (1) the baseline model, (2) a model with no unemployment insurance, and (3) a model in which households pool their idiosyncratic unemployment risk perfectly. In this version with pooled unemployment risk, I assume that there is no unemployment insurance. Thus, all households receive labor income of $N_t w_t z(1 - \tau_l)$, regardless of whether or not they are working. Note, in this model, households still face income risk due to idiosyncratic productivity shocks. They simply do not face additional income risk due to unemployment.

In these alternate versions of the model, I adjust \bar{w} so that the unemployment rate remains at 6% in the original steady-state. I also assume that the steady-state real interest rate remains at 1% in each version of the model. Table A.1 in the Online Appendix shows that the steady-state wealth distributions are similar in all three versions of the model, highlighting the relatively limited role that unemployment risk plays in determining the steady-state wealth distribution. In response to the aggregate shock, I assume that government spending adjusts to balance the government's budget constraint each period. In the Online Appendix, I show that results are similar with an alternative fiscal policy in which the lump-sum transfer adjusts.

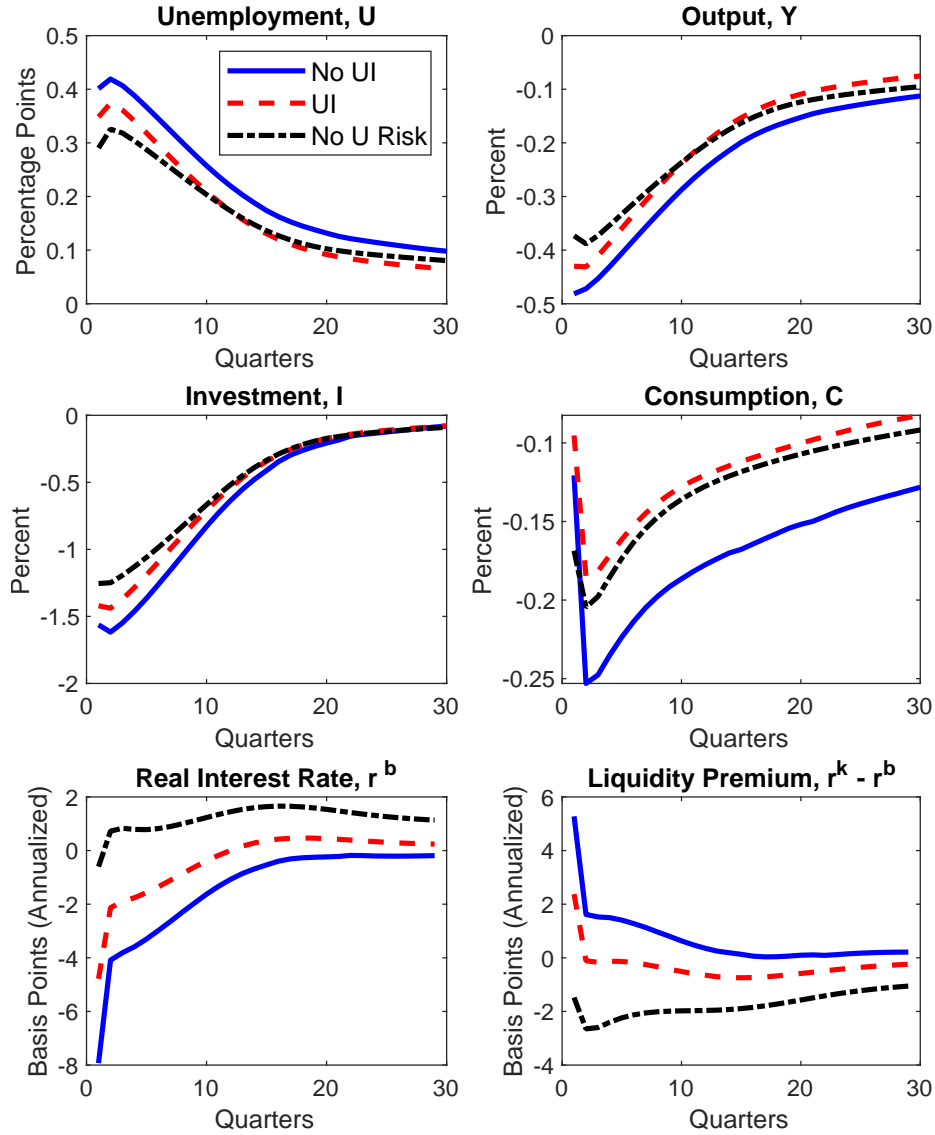
By comparing the response to the shock in the model with no unemployment insurance and in the model with pooled unemployment risk, I am able to assess the impact of unemployment risk on aggregate fluctuations. Between these extreme scenarios, the baseline model then shows the degree to which unemployment insurance is able to mitigate any amplification due to idiosyncratic unemployment risk. Figure 5 plots the impulse response of key variables to the aggregate productivity shock in each version of the model.

In all versions of the model, the decline in aggregate productivity causes a decline in vacancy posting and a rise in the unemployment rate. In response to an increase in unemployment risk, there is a flight-to-liquidity: demand for liquid assets increases, as these are best-suited to smoothing consumption during unemployment spells. Investment in capital falls, as employed households postpone investing in illiquid assets, and unemployed households withdraw from their illiquid asset holdings. In the presence of nominal rigidities, this decline in investment demand lowers aggregate output, raises unemployment, and initiates a feedback loop between unemployment risk and aggregate demand in the economy.

This mechanism is not operative if unemployment risk is pooled, and it is dampened if households have access to unemployment insurance. By providing a source of income during unemployment spells, unemployment insurance lessens the need

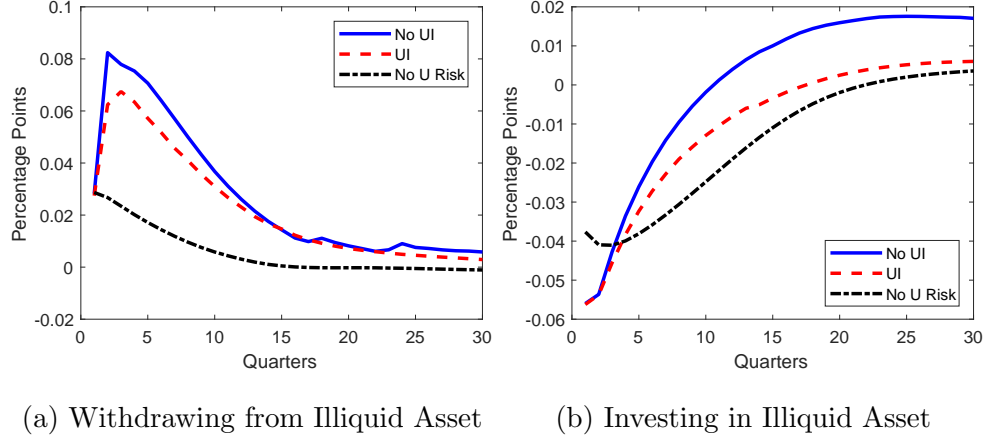
²⁸I consider a shock which lowers aggregate productivity by 0.1% on impact, and has a quarterly persistence equal to 0.9.

Figure 5. : Response to an Aggregate Productivity Shock



Notes: This figure shows the response of the economy to a negative aggregate productivity shock. "No UI" refers to the model without unemployment insurance. "UI" refers to the baseline model with unemployment insurance. "No U Risk" refers to the model in which households pool unemployment risk. See the text for full details.

Figure 6. : Response of Illiquid Asset Adjustment Probabilities



Notes: This figure shows the response of illiquid asset adjustment probabilities following a negative aggregate productivity shock. “No UI” refers to the model without unemployment insurance. “UI” refers to the baseline model with unemployment insurance. “No U Risk” refers to the model in which households pool unemployment risk. See the text for full details.

for holding liquid assets to smooth consumption during such times. Figure 6 shows the response of illiquid asset adjustment probabilities in each version of the model. The key difference is that without unemployment insurance there is a much larger increase in the probability of illiquid asset withdrawal in comparison to the other two versions of the model.

The quantitative significance of this mechanism can be seen in Figure 5. The main result is that the unemployment rate rises by around 30% more in the version without unemployment insurance than in the version with no unemployment risk, and that unemployment insurance removes around half of this amplification. The more unemployment risk that households face, the larger is the decline in investment, and the sharper is the decline in the real interest rate. The bottom-right panel of Figure 5 plots the liquidity premium, the spread between the rate of return on capital and the real interest rate. This spread is counter-cyclical when households face unemployment risk, and particularly so if there is no unemployment insurance.

A. Robustness

In Section J of the Online Appendix, I undertake a number of robustness exercises. The amplification implied by unemployment risk is robust to heterogeneous job separation rates, assuming that transfers rather than government spending adjust to balance the government’s budget constraint, a Taylor rule featuring interest-rate smoothing, assuming that profits are distributed directly to households, a wide range of values for the responsiveness of wages, and an alternative wage

rule in which wages respond to unemployment risk.²⁹ I also show that amplification is not confined to aggregate productivity shocks: I find similar amplification when studying a shock to the marginal efficiency of investment. This is significant, as such shocks have been shown to play an important role in business-cycle fluctuations in estimated DSGE models.

B. The Endogenous Response of Income Risk


In this section, I show that the endogenous response of income risk to the aggregate shock in the model is consistent with empirical evidence from the CPS. Guvenen, Ozkan and Song (2014) use Social Security Administration data to show that the skewness of the income growth distribution is strongly pro-cyclical: recessions are times when large negative income changes become much more likely. Using data from the March supplement of the CPS, I am able to break down income growth into hours growth and wage growth. Figure 7 shows that the pro-cyclical skewness of income growth is entirely driven by the pro-cyclical skewness of hours growth, while the distribution of hourly wage growth doesn't vary over the business cycle. Thus, large negative income changes in recessions become more likely due to an increased likelihood of a large decline in hours worked, i.e. an unemployment spell. In the Online Appendix, I provide more detail on the CPS data and additional evidence that the cyclicalities of income growth is driven by the extensive margin, specifically the cyclicalities of unemployment risk.

Figure 8 shows the effect of the aggregate shock on the skewness of the hours growth, wage growth and income growth distributions in the model. As in the data, the skewness of income growth is pro-cyclical, and it is driven entirely by the skewness of hours growth, which is around twice as volatile as the skewness of income growth. In the model, the skewness of the wage growth distribution is acyclical by construction, as it depends only on the exogenous stochastic process for idiosyncratic productivity.

Figure 8 is also useful for understanding why the flight-to-liquidity mechanism is not operative in the version of the model where unemployment risk is pooled. In this version of the model, the only source of income risk comes from idiosyncratic productivity shocks, so the skewness of the income growth distribution is unaffected by changes in the job-finding rate.


²⁹It would be interesting to extend the model to allow for varying search effort on the part of unemployed agents, for an hours choice on the intensive margin. I have not undertaken these extensions due to the computational complexity of the current model without these extra decisions on behalf of households. The same is true of an extension to accurately capture the earnings cost of job loss. Huckfeldt (2022) shows that such a cost is concentrated among individuals who change occupation. I believe that studying how such occupation choice interacts with liquid and illiquid asset holdings is an interesting topic for future research.

Figure 7. : Breakdown of Income Growth Skewness in the Data



Notes: Data from the Current Population Survey. Skewness measured using Pearson's second skewness coefficient (median skewness). Gaps in the data occur due to periods in which it is not possible to link individuals in the CPS over time.

Figure 8. : Model Response of Income Risk



Notes: This figure shows the response of income risk to a negative productivity shock. "No UI" refers to the model without unemployment insurance. "UI" refers to the baseline model with unemployment insurance. "No U Risk" refers to the model in which households pool unemployment risk. See the text for full details. Skewness measured using Pearson's second skewness coefficient (median skewness).

C. The Importance of Unemployment Insurance at the ZLB

I now consider how the importance of unemployment insurance as an automatic stabilizer depends on the responsiveness of monetary policy. I consider the response of the economy to the same aggregate productivity shock considered previously. However, I now assume that there is an exogenous lower bound on the nominal interest rate, such that monetary policy follows a truncated Taylor rule:

$$(32) \quad i_t = \max\{\bar{r}^b + \psi \log(\Pi_t), \underline{i}\}$$

I set \underline{i} such that monetary policy is constrained for 2 quarters in the baseline version of the model.³⁰ Figure 9 compares the impulse response functions of the baseline model with those from the versions of the model without unemployment insurance or without unemployment risk.

When monetary policy is constrained, the decline in investment demand that follows the increase in unemployment risk is not offset by lower interest rates. This strengthens the feedback loop between aggregate demand and unemployment risk, and increases the amplification coming from the flight-to-liquidity mechanism. Unemployment insurance plays a much more important role than in normal times: without unemployment insurance, unemployment rises by more than twice as much as in the baseline model. Monetary policy is constrained for longer, and both investment and inflation decline significantly more than with the baseline level of unemployment insurance.

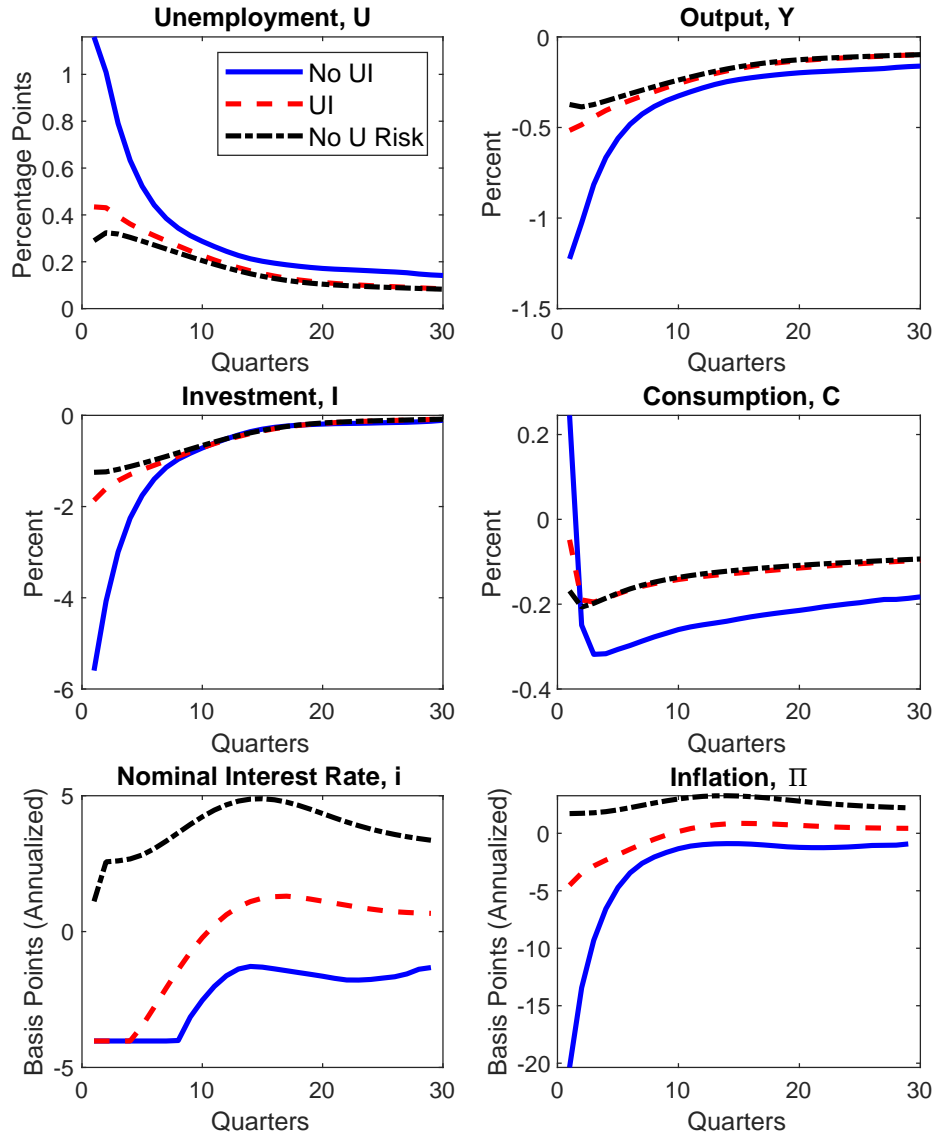
D. A Comparison with One-Asset Models

The previous sections have provided a quantitative assessment of the amplification provided by unemployment risk in a model which matches both liquid and illiquid wealth distributions. In order to understand the source of this amplification, I now briefly compare the results from Section VII with those from various models in which households only trade one (liquid) asset. The full results are provided in Section H of the Online Appendix.

In the first one-asset model, I remove bonds from the model and assume that households trade capital with no adjustment costs. Aside from these changes, I keep the parameterization as described in Section V.A. In such a model, households are well insured against unemployment risk, and consequently the decline in consumption during unemployment is smaller than documented in Section II. When studying the response of the economy to the aggregate productivity shock

³⁰The standard method for engineering a ZLB episode in New Keynesian models is a temporary rise in the discount factor, β . This does not work in this model due to the presence of capital and labor market frictions. Increasing the discount factor leads to a decline in unemployment, both because of an increase in the capital stock, which increases labor productivity, but also because a higher discount factor raises the value of a filled vacancy to the labor agency.

Figure 9. : Response to Shock with Constrained Monetary Policy



Notes: This figure shows the response of the economy to a negative productivity shock in the presence of a lower bound on the nominal interest rate. “No UI” refers to the model without unemployment insurance. “UI” refers to the baseline model with unemployment insurance. “No U Risk” refers to the model in which households pool unemployment risk. See the text for full details.

in this model, I find no amplification from unemployment risk. My findings are consistent with Gornemann, Kuester and Nakajima (2016). In such a model, while it increases the volatility of consumption, unemployment risk decreases the volatility of investment, and leaves the volatility of output or unemployment broadly unchanged.³¹

In the second one-asset model, I keep the liquid capital framework but lower the calibrated discount factor, β , from 0.982 to 0.96, in order to match the estimated consumption decline during unemployment. Here I find that unemployment risk actually *dampens* business cycle fluctuations slightly. A rise in unemployment risk leads to an increase in precautionary saving in capital. This dampens the decline in investment sufficiently such that output and employment actually fall less with unemployment risk than without. In the language of Challe et al. (2017), the stabilizing “aggregate supply” channel of unemployment risk dominates in this setting.

Finally, I also consider a one-asset model in which I remove capital. Thus, households only trade liquid bonds, and the potential stabilizing effect of unemployment risk operating through capital is absent. As in the models with liquid capital, I stick as close as possible to the calibration of the two-asset model. Here I find that, if the discount factor, β , is calibrated to match the consumption decline during unemployment seen in the data, there is only a small amount of amplification from unemployment risk.

Overall, studying these alternative one-asset models suggests that the inclusion of illiquid wealth and the flight-to-liquidity motive is key in generating the amplification seen in the two-asset model in Section VII.

VIII. Temporary Tax-Free Illiquid Asset Withdrawals

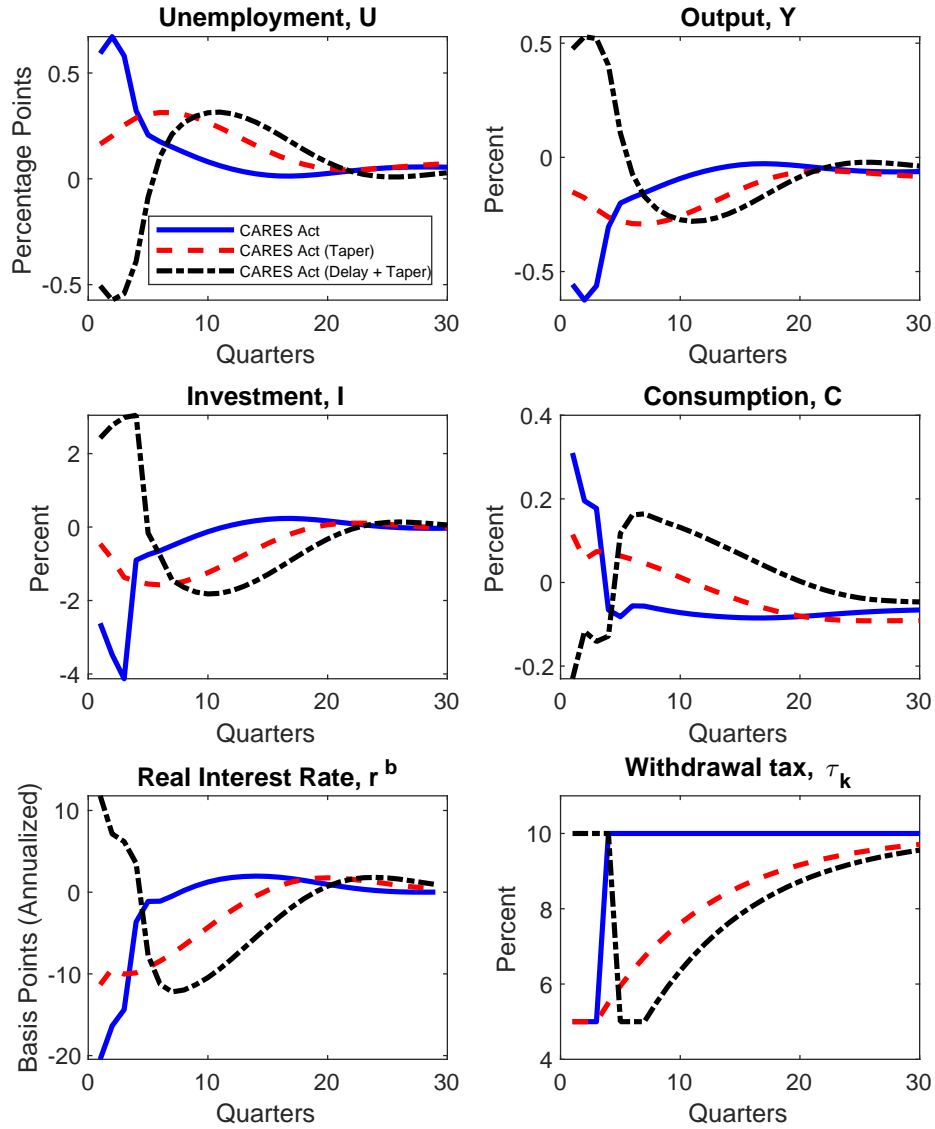
In response to the COVID-19 pandemic, the CARES Act removed the 10% early withdrawal penalty on retirement accounts until the end of 2020 for individuals that experience adverse financial consequences due to the pandemic. The two-asset model in this paper is a useful laboratory for studying the effect of such a policy.³²

In this section, I investigate the implications of variations of this policy in the baseline calibration of the model. I consider three alternative implementations. First, and most similar to the CARES Act, I assume that the withdrawal tax on illiquid assets is reduced unexpectedly and immediately from 10% to 5% for three quarters for individuals that are unemployed. It then returns immediately to 10%. I do not remove the tax entirely as in reality retirement accounts are only a

³¹These results also consistent with Cho (2023) who shows that time-varying unemployment risk has little effect on business cycle volatility in an estimated HANK model without illiquid assets.

³²Bayer et al. (2023) also study the CARES Act in a two-asset HANK model. They do not consider the effects of the removal of the early withdrawal penalty, instead focusing on the stimulus provided by the large transfer payments in the bill.

Figure 10. : Response to Temporary Reductions of Withdrawal Tax



Notes: This figure shows the response to various temporary reductions of the illiquid asset withdrawal tax. The path of the withdrawal tax in each case is shown in the bottom-right panel. See the text for full details.

fraction of total illiquid asset holdings. Second, I consider a variation where after three quarters the withdrawal tax returns to 10% much more slowly.³³ Third, I consider a variation where the reduction in the withdrawal tax is pre-announced by one year. As in the second case, the tax is then cut to 5% for three quarters and returns to 10% only gradually. For simplicity, I assume that these tax cuts occur in the steady-state of the model, rather than layering them on top of an additional aggregate shock to mimic the COVID-19 recession.

Figure 10 shows the results of these policies. The three alternate paths of the withdrawal tax are shown in the bottom-right panel. Starting with the first scenario, labelled “CARES Act”, in which the tax is cut immediately and then returns to 10% after three quarters, we see that while the policy helps unemployed individuals smooth their consumption, it actually has a sharp contractionary effect. The policy leads to a synchronized withdrawal of illiquid assets, and consequently a significant decline in investment in capital. This is only partially offset by an increase in consumption. The overall effect is a decline in aggregate demand, a fall in the equilibrium interest rate, and consequently a rise in unemployment.

By considering the second case, labelled “CARES Act (Taper)”, we see that allowing the tax cut to last longer significantly lowers the peak rise in unemployment, or the decline in output on impact. This occurs as such an implementation avoids the synchronized withdrawal of illiquid assets that occurs when the tax cut is known to only last for three quarters.

Finally, the third case, labelled “CARES Act (Delay + Taper)” shows that both pre-announcing the policy and tapering its withdrawal actually leads to an expansion in the short-run. Knowing that the withdrawal tax will be lower in the future, households significantly increase their investment in illiquid assets in the first year, causing an expansion in output and a reduction in unemployment.

IX. Conclusion

This paper shows that the combination of endogenous unemployment risk and the presence of illiquid assets provides an important propagation mechanism for aggregate shocks: higher unemployment risk leads to a flight-to-liquidity and initiates a feedback loop between unemployment risk and aggregate demand. Unemployment insurance plays an important role as an automatic stabilizer, particularly if monetary policy is constrained.

The two-asset model is consistent with new empirical evidence on the relationship between unemployment, the liquidity of asset holdings and consumption. Using data from the Consumer Expenditure Survey, I find that the consumption decline during unemployment is largest for poor hand-to-mouth households, smaller for the wealthy hand-to-mouth, and smallest for the non hand-to-mouth. The two-

³³Specifically, following an AR(1) process with a quarterly persistence of 0.9.

asset model is able to match this finding due to the costs associated with adjusting illiquid asset holdings. Some wealthy hand-to-mouth households pay these adjustment costs, and consequently are able to smooth their consumption as well as the non hand-to-mouth, while others do not pay the adjustment costs and are unable to smooth their consumption, like poor hand-to-mouth households.

In the model, unemployed households do not need to withdraw from their illiquid asset holdings until they have first run down their liquid asset holdings. However, when their liquid asset holdings are depleted, they are then likely to withdraw from their illiquid asset holdings. Consequently, unemployed households are more likely to make a withdrawal from their illiquid asset holdings than employed households, particularly if their unemployment spell is long or their liquid asset holdings are low. Using data from the Survey of Consumer Finances, I show that these patterns are confirmed in the data.

I also provide evidence in favor of the flight-to-liquidity mechanism that is central to the model's response to aggregate shocks: using data from the Survey of Consumer Expectations I show that consumption declines when idiosyncratic unemployment risk rises, consistent with a rise in precautionary saving, but that investment in illiquid assets declines, consistent with a shift in portfolios towards liquid assets. This pattern also occurs in the model.

The model suggests that an important role for unemployment insurance is its ability to dampen aggregate fluctuations by lessening the flight-to-liquidity that occurs when unemployment risk is heightened. However, the model has abstracted from search effort on the part of unemployed workers, or a mechanism by which unemployment insurance affects the level of wages. Such features would imply that there is an important trade-off between the effect of unemployment insurance on the volatility of the unemployment rate and the effect on the average level of unemployment. I leave an investigation of this trade-off to future work.

REFERENCES

- Aguiar, Mark, and Erik Hurst.** 2005. "Consumption versus Expenditure." *Journal of Political Economy*, 113(5): 919–948.
- Basten, Christoph, Andreas Fagereng, and Kjetil Telle.** 2016. "Saving and portfolio allocation before and after job loss." *Journal of Money, Credit and Banking*, 48(2-3): 293–324.
- Bayer, Christian, Benjamin Born, Ralph Luetticke, and Gernot J Müller.** 2023. "The coronavirus stimulus package: How large is the transfer multiplier." *The Economic Journal*, 133(652): 1318–1347.
- Bayer, Christian, Ralph Lütticke, Lien Pham-Dao, and Volker Tjaden.** 2019. "Precautionary savings, illiquid assets, and the aggregate consequences of shocks to household income risk." *Econometrica*, 87(1): 255–290.
- Bhutta, Neil, and Benjamin J Keys.** 2016. "Interest rates and equity extraction during the housing boom." *American Economic Review*, 106(7): 1742–74.
- Challe, Edouard, Julien Matheron, Xavier Ragot, and Juan F Rubio-Ramirez.** 2017. "Precautionary saving and aggregate demand." *Quantitative Economics*, 8(2): 435–478.
- Cho, Daeha.** 2023. "Unemployment risk, MPC heterogeneity, and business cycles." *Quantitative Economics*.
- Chodorow-Reich, Gabriel, and Loukas Karabarbounis.** 2016. "The cyclicity of the opportunity cost of employment." *Journal of Political Economy*, 124(6): 1563–1618.
- Christiano, Lawrence J, Martin S Eichenbaum, and Mathias Trabandt.** 2016. "Unemployment and business cycles." *Econometrica*, 84(4): 1523–1569.
- Coyne, David, Itzik Fadlon, and Tommaso Porzio.** 2022. "Measuring Valuation of Liquidity with Penalized Withdrawals." National Bureau of Economic Research.
- Den Haan, Wouter J, Garey Ramey, and Joel Watson.** 2000. "Job destruction and propagation of shocks." *American economic review*, 90(3): 482–498.
- Den Haan, Wouter J, Pontus Rendahl, and Markus Riegler.** 2017. "Unemployment (fears) and deflationary spirals." *Journal of the European Economic Association*, 16(5): 1281–1349.
- Department of Labor.** 2018. "Comparison of State Unemployment Laws."
- Gornemann, Nils, Keith Kuester, and Makoto Nakajima.** 2016. "Doves for the rich, hawks for the poor? Distributional consequences of monetary policy."

- Graves, Sebastian.** 2020. “Does Unemployment Risk Affect Business Cycle Dynamics?” *International Finance Discussion Paper*, , (1298).
- Greenwood, Jeremy, Zvi Hercowitz, and Gregory W Huffman.** 1988. “Investment, capacity utilization, and the real business cycle.” *The American Economic Review*, 402–417.
- Gruber, Jonathan.** 1997. “The Consumption Smoothing Benefits of Unemployment Insurance.” *The American Economic Review*, 87(1): 192–205.
- Güvenen, Fatih, Fatih Karahan, Serdar Ozkan, and Jae Song.** 2015. “What do data on millions of US workers reveal about life-cycle earnings risk?” National Bureau of Economic Research.
- Güvenen, Fatih, Serdar Ozkan, and Jae Song.** 2014. “The nature of countercyclical income risk.” *Journal of Political Economy*, 122(3): 621–660.
- Hagedorn, Marcus, and Iourii Manovskii.** 2008. “The cyclical behavior of equilibrium unemployment and vacancies revisited.” *American Economic Review*, 98(4): 1692–1706.
- Hoffmann, Eran B., and Davide Malacrino.** 2019. “Employment time and the cyclical growth of earnings.” *Journal of Public Economics*, 169(C): 160–171.
- Huckfeldt, Christopher.** 2022. “Understanding the scarring effect of recessions.” *American Economic Review*, 112(4): 1273–1310.
- Kaplan, Greg, Benjamin Moll, and Giovanni L Violante.** 2018. “Monetary policy according to HANK.” *American Economic Review*, 108(3): 697–743.
- Kaplan, Greg, Giovanni L. Violante, and Justin Weidner.** 2014. “The Wealthy Hand-to-Mouth.” *Brookings Papers on Economic Activity*, 45(1 (Spring)): 77–153.
- Kekre, Rohan.** 2016. “Unemployment Insurance in Macroeconomic Stabilization.” Society for Economic Dynamics 2016 Meeting Papers 1184.
- Khan, Aubhik, and Julia K Thomas.** 2008. “Idiosyncratic shocks and the role of nonconvexities in plant and aggregate investment dynamics.” *Econometrica*, 76(2): 395–436.
- Kolsrud, Jonas, Camille Landais, Peter Nilsson, and Johannes Spinnewijn.** 2018. “The optimal timing of unemployment benefits: Theory and evidence from Sweden.” *American Economic Review*, 108(4-5): 985–1033.
- McKay, Alisdair, and Ricardo Reis.** 2016. “The role of automatic stabilizers in the US business cycle.” *Econometrica*, 84(1): 141–194.

- McKay, Alisdair, and Ricardo Reis.** 2021. “Optimal automatic stabilizers.” *The Review of Economic Studies*, 88(5): 2375–2406.
- Ravn, Morten O., and Vincent Sterk.** 2017. “Job uncertainty and deep recessions.” *Journal of Monetary Economics*, 90(C): 125–141.
- Rotemberg, Julio J.** 1982. “Sticky prices in the United States.” *Journal of Political Economy*, 90(6): 1187–1211.